

# EFFICIENT CALIBRATION OF EMBEDDED MPC

Marco Forgone<sup>1</sup>, Dario Piga<sup>1</sup>, and Alberto Bemporad<sup>2</sup>

<sup>1</sup>IDSIA Dalle Molle Institute for Artificial Intelligence SUPSI-USI, Lugano, Switzerland

<sup>2</sup>IMT School for Advanced Studies Lucca, Lucca, Italy

IFAC 2020 World Congress  
Berlin, Germany

# Motivations

Calibration of **model predictive controllers** is **costly** and **time-consuming**. One has to choose model, cost function weights, prediction/control horizon, design a state estimator, etc.


The design is even more involved for **embedded applications**, where the hardware is also a limit. Need to choose carefully:

- Solution strategy (QP solver?)
- Low-level solver settings (e.g., tolerances)
- A **feasible** MPC sampling time

In this paper, we introduce an approach to ease embedded MPC calibration based on **global optimization** and repeated experiments.

# Control architecture

## Model Predictive Controller



Linear constrained MPC for a (non-linear) plant with model

$$\begin{aligned} \dot{x}(t) &= f(x(t), u(t); \beta) & x_{t+1} &= A(T_s^{\text{MPC}}, \beta)x_t + B(T_s^{\text{MPC}}, \beta)u_t \\ y(t) &= g(x(t); \beta) & \Rightarrow & y_t = C(\beta)x_t \end{aligned}$$

input  $u$  chosen on-line as the solution (applied in receding horizon) of:

$$\min_{\{u_{t+k|t}\}_{k=1}^{N_c}} \sum_{k=1}^{N_p} \left\| \hat{y}_{t+k|t} - \bar{y}_{t+k} \right\|_{Q_y}^2 + \left\| u_{t+k|t} - u_{t+k-1|t} \right\|_{Q_{\Delta u}}^2$$


s.t. model equations + constraints on  $y$ ,  $u$ ,  $\Delta u$

# Control architecture

## Design variables

The MPC design variables to be tuned are

- Model parameters:  $\beta$
- MPC sampling time:  $T_s^{\text{MPC}}$
- MPC cost function:  $Q_y, Q_{\Delta u}, N_c, N_p$
- MPC implementation: QP solver tolerances  $\epsilon_{\text{abs}}, \epsilon_{\text{rel}}$
- State estimator: Kalman covariances  $W_v, W_w$



# MPC calibration procedure

## Overview

In our optimization-based MPC calibration, we define

- Tunable **design parameters** collected in a **design vector**  $\theta \in \Theta$ .
- A closed-loop **performance index**  $J$  defined in terms of measured input/outputs during the calibration experiment:  $J = J(y_{1:T}, u_{1:T}; \theta)$
- An **procedure** to perform **calibration experiments** (or SIL simulations) representative of the intended closed-loop operation

MPC calibration is seen as a **global optimization** problem:

$$\begin{aligned} \theta^{\text{opt}} &= \arg \min_{\theta \in \Theta} J(y_{1:T}, u_{1:T}; \theta) \\ \text{s.t. } T_{\text{calc}}^{\text{MPC}}(\theta) &\leq T_s^{\text{MPC}} \end{aligned}$$

each (noisy) **function evaluation** correspond to a **calibration experiment**.  
Both  $J$  and  $T_{\text{calc}}^{\text{MPC}}$  (worst-case MPC computation time) are observed!

# MPC calibration procedure

## Overview

In our optimization-based MPC calibration, we define

- Tunable **design parameters** collected in a **design vector**  $\theta \in \Theta$ .
- A closed-loop **performance index**  $J$  defined in terms of measured input/outputs during the calibration experiment:  $J = J(y_{1:T}, u_{1:T}; \theta)$
- An **procedure** to perform **calibration experiments** (or SIL simulations) representative of the intended closed-loop operation

MPC calibration is seen as a **global optimization** problem:

$$\begin{aligned} \theta^{\text{opt}} &= \arg \min_{\theta \in \Theta} J(y_{1:T}, u_{1:T}; \theta) \\ \text{s.t. } T_{\text{calc}}^{\text{MPC}}(\theta) &\leq T_s^{\text{MPC}} \end{aligned}$$

each (noisy) **function evaluation** correspond to a **calibration experiment**. Both  $J$  and  $T_{\text{calc}}^{\text{MPC}}$  (worst-case MPC computation time) are observed!

# Global Optimization Algorithm

## Overview

Several **derivative-free** global optimization algorithms exist:

- Response surface methods
- Bayesian Optimization
- Genetic algorithms
- Particle Swarm Optimization
- ...

Here, we adopt **GLIS**: a method recently introduced in (Bemporad, 2019).

- Radial basis function **surrogate** + inverse distance weighting (IDW) for **exploration**
- Performs better than BO on standard benchmarks, at a lower computational cost

# Global Optimization Algorithm

## Overview

Several **derivative-free** global optimization algorithms exist:

- Response surface methods
- Bayesian Optimization
- Genetic algorithms
- Particle Swarm Optimization
- ...

Here, we adopt **GLIS**: a method recently introduced in (Bemporad, 2019).

- Radial basis function **surrogate** + inverse distance weighting (IDW) for **exploration**
- Performs better than BO on standard benchmarks, at a lower computational cost



# GLIS Algorithm

## Surrogate function

Goal: solve global optimization problem:


$$\min_{\theta \in \Theta} J(\theta)$$

- 1 Collect  $n_{\text{in}}$  initial observations  $\{(\theta_1, J_1), (\theta_2, J_2), \dots, (\theta_{n_{\text{in}}}, J_{n_{\text{in}}})\}$
- 2 Build a surrogate function

$$\hat{J}(\theta) = \sum_{i=1}^n \alpha_i \phi(\|\theta - \theta_i\|_2)$$

Minimizing  $\hat{J}(\theta)$  to find  $\theta_{n+1}$  may easily miss the global optimum!

True  $J(\theta)$ , surrogate  $\hat{J}(\theta)$



$\phi$  = radial basis function

Example:  $\phi(d) = \frac{1}{1+(\epsilon d)^2}$

# GLIS Algorithm

## Surrogate function

Goal: solve global optimization problem:


$$\min_{\theta \in \Theta} J(\theta)$$

- 1 Collect  $n_{\text{in}}$  initial observations  $\{(\theta_1, J_1), (\theta_2, J_2), \dots, (\theta_{n_{\text{in}}}, J_{n_{\text{in}}})\}$
- 2 Build a surrogate function

$$\hat{J}(\theta) = \sum_{i=1}^n \alpha_i \phi(\|\theta - \theta_i\|_2)$$

Minimizing  $\hat{J}(\theta)$  to find  $\theta_{n+1}$  may easily miss the global optimum!

True  $J(\theta)$ , surrogate  $\hat{J}(\theta)$



$\phi$  = radial basis function

Example:  $\phi(d) = \frac{1}{1+(\epsilon d)^2}$

# GLIS Algorithm

## Exploration vs. Exploitation

- 3 Construct the IDW exploration function

$$z(\theta) = \frac{2}{\pi} \Delta F \tan^{-1} \left( \frac{1}{\sum_{i=1}^n w_i(\theta)} \right)$$

$$\text{where } w_i(\theta) = \frac{e^{-\|\theta - \theta_i\|^2}}{\|\theta - \theta_i\|^2}$$


- 4 Optimize the acquisition function:

$$\theta_{n+1} = \arg \min_{\theta \in \Theta} \hat{J}(\theta) - \delta z(\theta)$$

to get the query point  $\theta_{n+1}$ . Execute experiment with  $\theta_{n+1}$ , measure  $J_{n+1}$ .

Iterate the procedure for  $n + 2, n + 3 \dots$

Exploration function  $z(\theta)$



$\delta$ : exploitation vs.  
exploration trade-off

# GLIS Algorithm

## Exploration vs. Exploitation

- 3 Construct the IDW exploration function

$$z(\theta) = \frac{2}{\pi} \Delta F \tan^{-1} \left( \frac{1}{\sum_{i=1}^n w_i(\theta)} \right)$$

$$\text{where } w_i(\theta) = \frac{e^{-\|\theta - \theta_i\|^2}}{\|\theta - \theta_i\|^2}$$


- 4 Optimize the acquisition function:

$$\theta_{n+1} = \arg \min_{\theta \in \Theta} \hat{J}(\theta) - \delta z(\theta)$$

to get the query point  $\theta_{n+1}$ . Execute experiment with  $\theta_{n+1}$ , measure  $J_{n+1}$ .

Iterate the procedure for  $n + 2, n + 3 \dots$


Exploration function  $z(\theta)$



$\delta$ : exploitation vs.  
exploration trade-off

# MPC Calibration Example

## Cart-pole system



- Input: force  $F$  with **fast disturbance** (5 rad/sec)
- Output: noisy position  $p$
- Objective: follow trajectory for  $p$ , keep  $\phi$  small.

## Optimization-based calibration of

- 1 MPC sample time  $T_s^{\text{MPC}}$
- 2 MPC weights  $Q_y$  and  $Q_{\Delta u}$
- 3 Prediction and control horizon
- 4 Kalman filter covariance matrices
- 5 QP solver's abs. and rel. tolerances


$n_\theta = 14$  design parameters optimized using  $n = 500$  iterations of GLIS

$$J = \int_0^{T_{\text{exp}}} |\bar{p}(t) - p(t)| + 30|\phi(t)| dt$$

Python codes on-line, 100% **open source** (including all dependencies).  
<https://github.com/forgi86/efficient-calibration-embedded-MPC>

# MPC Calibration Example

## Cart-pole system



- Input: force  $F$  with **fast disturbance** (5 rad/sec)
- Output: noisy position  $p$
- Objective: follow trajectory for  $p$ , keep  $\phi$  small.

## Optimization-based calibration of


- 1 MPC sample time  $T_s^{\text{MPC}}$
- 2 MPC weights  $Q_y$  and  $Q_{\Delta u}$
- 3 Prediction and control horizon
- 4 Kalman filter covariance matrices
- 5 QP solver's abs. and rel. tolerances

$n_\theta = 14$  design parameters optimized using  $n = 500$  iterations of GLIS


$$J = \int_0^{T_{\text{exp}}} |\bar{p}(t) - p(t)| + 30|\phi(t)| dt$$

Python codes on-line, 100% **open source** (including all dependencies).  
<https://github.com/forgi86/efficient-calibration-embedded-MPC>

# MPC Calibration Example




VS.



- An Intel i5 PC (left) vs. an ARM-based Raspberry Pi 3 (right)
- PI is about 10 times slower than the PC for MPC computations


# MPC Calibration Example

## Optimized MPC on the PC



$$T_s^{\text{MPC}} = 6 \text{ ms}$$

## Optimized MPC on the Raspberry Pi




$$T_s^{\text{MPC}} = 22 \text{ ms}$$

- Position and angle control tighter on the PC
- Faster loop update on the PC  $\Rightarrow$  more effective disturbance rejection
- Calibration squeezes max performance out of the hardware!




# MPC Calibration Example

## GLIS iterations on the PC



## GLIS iterations on the Raspberry PI




- High cost for interrupted tests/unfeasible design
- For increasing  $n$ , more points have “low” cost
- Experiments with high cost still appear for large  $n$  (exploration)

# Simulation Example


## Cart-pole system

Similar calibration results are obtained using Bayesian Optimization (BO). However, BO is a much heavier algorithm!

### GLIS iterations on the PI



### BO iterations on the PI



- With GLIS, the overall computation time is dominated by function evaluations (= running MPC and simulating the system).

# Conclusions

An approach for embedded MPC calibration approach based on global optimization and repeated experiments

- Higher- and lower-level tuning knobs jointly optimized for performance, while keeping the design feasible for real-time implementation.
- Tested in simulation on two hardware platforms

Current/future works

- Application to robotic systems
- Preference-based MPC calibration

# Conclusions

An approach for embedded MPC calibration approach based on global optimization and repeated experiments

- Higher- and lower-level tuning knobs jointly optimized for performance, while keeping the design feasible for real-time implementation.
- Tested in simulation on two hardware platforms

Current/future works

- Application to robotic systems
- Preference-based MPC calibration

Thank you.  
Questions?

`marco.forgione@idsia.ch`